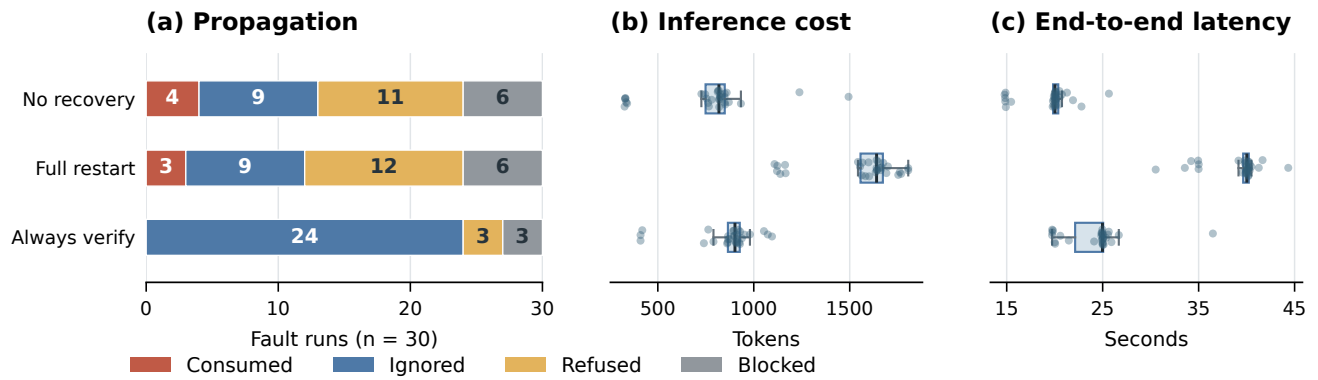
(a) Propagation
No recovery
Full restart
Always verify
4
9
11
6
3
9
12
6
24
3
3
0
10
20
30
Fault runs (n = 30)
Consumed
Ignored
Refused
Blocked
(b) Inference cost
500
1000
1500
Tokens
(c) End-to-end latency
15
25
35
45
Seconds