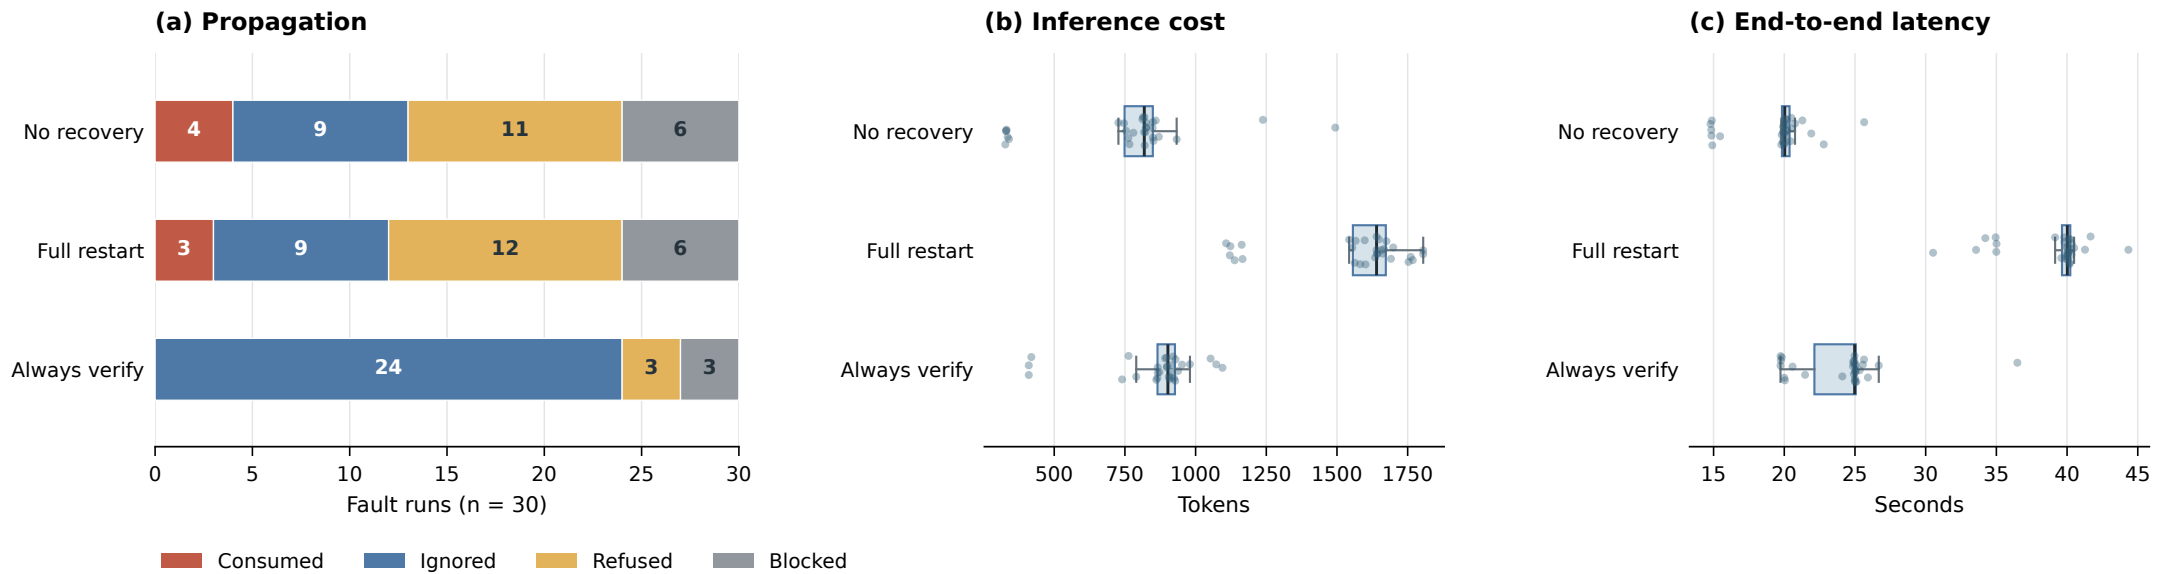
(a) Propagation
No recovery
Full restart
Always verify
4
9
11
6
3
9
12
6
24
3
3
0
5
10
15
20
25
30
Fault runs (n = 30)
Consumed
Ignored
Refused
Blocked
(b) Inference cost
No recovery
Full restart
Always verify
500
750
1000
1250
1500
1750
Tokens
(c) End-to-end latency
No recovery
Full restart
Always verify
15
20
25
30
35
40
45
Seconds